Expert Oracle Database Architecture

9*i* and 10*g* Programming Techniques and Solutions

Thomas Kyte

Apress®

Expert Oracle Database Architecture: 9*i* and 10*g* Programming Techniques and Solutions

Copyright © 2005 by Thomas Kyte

All rights reserved. No part of this work may be reproduced or transmitted in any form or by any means, electronic or mechanical, including photocopying, recording, or by any information storage or retrieval system, without the prior written permission of the copyright owner and the publisher.

ISBN (pbk): 1-59059-530-0

Printed and bound in the United States of America 9 8 7 6 5 4 3 2 1

Trademarked names may appear in this book. Rather than use a trademark symbol with every occurrence of a trademarked name, we use the names only in an editorial fashion and to the benefit of the trademark owner, with no intention of infringement of the trademark.

Lead Editor: Tony Davis
Technical Reviewer: Jonathan Lewis, Roderick Manalac, Michael Möller, Gabe Romanescu
Editorial Board: Steve Anglin, Dan Appleman, Ewan Buckingham, Gary Cornell, Tony Davis, Jason Gilmore, Jonathan Hassell, Chris Mills, Dominic Shakeshaft, Jim Sumser
Associate Publisher: Grace Wong
Project Manager: Sofia Marchant
Copy Edit Manager: Nicole LeClerc
Assistant Production Director: Kari Brooks-Copony
Production Editor: Katie Stence
Compositor: Dina Quan
Proofreader: Linda Marousek
Indexer: Broccoli Information Management
Artist: Kinetic Publishing Services, LLC
Interior Designer: Van Winkle Design Group
Cover Designer: Kurt Krames
Manufacturing Director: Tom Debolski

Distributed to the book trade worldwide by Springer-Verlag New York, Inc., 233 Spring Street, 6th Floor, New York, NY 10013. Phone 1-800-SPRINGER, fax 201-348-4505, e-mail `orders-ny@springer-sbm.com`, or visit `http://www.springeronline.com`.

For information on translations, please contact Apress directly at 2560 Ninth Street, Suite 219, Berkeley, CA 94710. Phone 510-549-5930, fax 510-549-5939, e-mail `info@apress.com`, or visit `http://www.apress.com`.

The information in this book is distributed on an "as is" basis, without warranty. Although every precaution has been taken in the preparation of this work, neither the author(s) nor Apress shall have any liability to any person or entity with respect to any loss or damage caused or alleged to be caused directly or indirectly by the information contained in this work.

The source code for this book is available to readers at `http://www.apress.com` in the Source Code section.

Contents

Foreword xi
About the Author xiv
About the Technical Reviewers xv
Acknowledgments xvi
Introduction xvii
Setting Up Your Environment xxv

CHAPTER 1 Developing Successful Oracle Applications 1

My Approach 2
The Black Box Approach 4
How (and How Not) to Develop Database Applications 9
Understanding Oracle Architecture 9
Understanding Concurrency Control 15
Multi-Versioning 20
Database Independence? 26
"How Do I Make It Run Faster?" 42
The DBA–Developer Relationship 46
Summary 47

CHAPTER 2 Architecture Overview 49

Defining Database and Instance 50
The SGA and Background Processes 55
Connecting to Oracle 57
Dedicated Server 57
Shared Server 59
Mechanics of Connecting over TCP/IP 60
Summary 62

CHAPTER 3 Files 65

Parameter Files 66
What Are Parameters? 67
Legacy init.ora Parameter Files 69
Server Parameter Files (SPFILEs) 71
Parameter File Wrap-Up 78
Trace Files 78
Requested Trace Files 79
Trace Files Generated in Response to Internal Errors 83
Trace File Wrap-Up 85
Alert File 85
Data Files 88
A Brief Review of File System Mechanisms 89
The Storage Hierarchy in an Oracle Database 90
Dictionary-Managed and Locally-Managed Tablespaces 94
Temp Files 96
Control Files 98
Redo Log Files 98
Online Redo Log 99
Archived Redo Log 101
Password Files 103
Change Tracking File 106
Flashback Log Files 107
Flashback Database 107
Flash Recovery Area 108
DMP Files (EXP/IMP Files) 108
Data Pump Files 110
Flat Files 113
Summary 114

CHAPTER 4 Memory Structures 115

The Process Global Area and User Global Area 115
Manual PGA Memory Management 116
Automatic PGA Memory Management 123
Choosing Between Manual and Auto Memory Management 133
PGA and UGA Wrap-Up 135
The System Global Area 135
Fixed SGA 139
Redo Buffer 140

Block Buffer Cache 141
Shared Pool 148
Large Pool 150
Java Pool 151
Streams Pool 152
Automatic SGA Memory Management 152
Summary 154

CHAPTER 5 Oracle Processes 155

Server Processes 156
Dedicated Server Connections 156
Shared Server Connections 158
Connections vs. Sessions 159
Dedicated Server vs. Shared Server 165
Dedicated/Shared Server Wrap-Up 169
Background Processes 170
Focused Background Processes 171
Utility Background Processes 178
Slave Processes 181
I/O Slaves 181
Parallel Query Slaves 182
Summary 182

CHAPTER 6 Locking and Latching 183

What Are Locks? 183
Locking Issues 186
Lost Updates 186
Pessimistic Locking 187
Optimistic Locking 189
Optimistic or Pessimistic Locking? 200
Blocking 200
Deadlocks 203
Lock Escalation 208
Lock Types 209
DML Locks 209
DDL Locks 217
Latches 220
Manual Locking and User-Defined Locks 229
Summary 230

CHAPTER 7 Concurrency and Multi-Versioning 231

What Are Concurrency Controls? 231
Transaction Isolation Levels 232
READ UNCOMMITTED 234
READ COMMITTED 235
REPEATABLE READ 237
SERIALIZABLE 239
READ ONLY 241
Implications of Multi-Version Read Consistency 242
A Common Data Warehousing Technique That Fails 242
An Explanation for Higher Than Expected I/O on Hot Tables 244
Write Consistency 246
Consistent Reads and Current Reads 247
Seeing a Restart 249
Why Is a Restart Important to Us? 252
Summary 253

CHAPTER 8 Transactions 255

Transaction Control Statements 256
Atomicity 257
Statement-Level Atomicity 257
Procedure-Level Atomicity 259
Transaction-Level Atomicity 262
Integrity Constraints and Transactions 262
IMMEDIATE Constraints 262
DEFERRABLE Constraints and Cascading Updates 263
Bad Transaction Habits 265
Committing in a Loop 266
Using Autocommit 272
Distributed Transactions 273
Autonomous Transactions 275
How Autonomous Transactions Work 275
When to Use Autonomous Transactions 277
Summary 281

CHAPTER 9 Redo and Undo 283

What Is Redo? 283
What Is Undo? 284
How Redo and Undo Work Together 287
Example INSERT-UPDATE-DELETE Scenario 287
Commit and Rollback Processing 291
What Does a COMMIT Do? 292
What Does a ROLLBACK Do? 298
Investigating Redo 300
Measuring Redo 300
Redo Generation and BEFORE/AFTER Triggers 302
Can I Turn Off Redo Log Generation? 308
Why Can't I Allocate a New Log? 313
Block Cleanout 314
Log Contention 317
Temporary Tables and Redo/Undo 319
Investigating Undo 323
What Generates the Most and Least Undo? 323
ORA-01555: snapshot too old Error 325
Summary 336

CHAPTER 10 Database Tables 337

Types of Tables 337
Terminology 339
Segment 339
Segment Space Management 341
High-Water Mark 342
FREELISTS 344
PCTFREE and PCTUSED 347
LOGGING and NOLOGGING 350
INITRANS and MAXTRANS 351
Heap Organized Tables 351
Index Organized Tables 354
Index Organized Tables Wrap-Up 369
Index Clustered Tables 370
Index Clustered Tables Wrap-Up 378
Hash Clustered Tables 378
Hash Clustered Tables Wrap-Up 387
Sorted Hash Clustered Tables 388

Nested Tables 390
Nested Tables Syntax 391
Nested Table Storage 399
Nested Tables Wrap-Up 402
Temporary Tables 403
Temporary Tables Wrap-Up 410
Object Tables 410
Object Tables Wrap-Up 418
Summary 418

CHAPTER 11 Indexes 421

An Overview of Oracle Indexes 422
B*Tree Indexes 423
Index Key Compression 426
Reverse Key Indexes 429
Descending Indexes 435
When Should You Use a B*Tree Index? 437
B*Trees Wrap-Up 447
Bitmap Indexes 448
When Should You Use a Bitmap Index? 449
Bitmap Join Indexes 453
Bitmap Indexes Wrap-Up 455
Function-Based Indexes 455
Important Implementation Details 455
A Simple Function-Based Index Example 456
Indexing Only Some of the Rows 464
Implementing Selective Uniqueness 466
Caveat on CASE 467
Caveat Regarding ORA-01743 468
Function-Based Indexes Wrap-Up 469
Application Domain Indexes 469
Frequently Asked Questions and Myths About Indexes 471
Do Indexes Work on Views? 471
Do Nulls and Indexes Work Together? 471
Should Foreign Keys Be Indexed? 474
Why Isn't My Index Getting Used? 475
Myth: Space Is Never Reused in an Index 482
Myth: Most Discriminating Elements Should Be First 485
Summary 488

CHAPTER 12 Datatypes 489

An Overview of Oracle Datatypes 489
Character and Binary String Types 492
NLS Overview 492
Character Strings 495
Binary Strings: RAW Types 502
Number Types 504
NUMBER Type Syntax and Usage 507
BINARY_FLOAT/BINARY_DOUBLE Type Syntax and Usage 510
Non-Native Number Types 511
Performance Considerations 511
LONG Types 513
Restrictions on LONG and LONG RAW Types 513
Coping with Legacy LONG Types 515
DATE, TIMESTAMP, and INTERVAL Types 520
Formats 521
DATE Type 522
TIMESTAMP Type 529
INTERVAL Type 537
LOB Types 540
Internal LOBs 541
BFILEs 553
ROWID/UROWID Types 555
Summary 556

CHAPTER 13 Partitioning 557

Partitioning Overview 558
Increased Availability 558
Reduced Administrative Burden 560
Enhanced Statement Performance 565
Table Partitioning Schemes 567
Range Partitioning 567
Hash Partitioning 570
List Partitioning 575
Composite Partitioning 577
Row Movement 579
Table Partitioning Schemes Wrap-Up 581

Partitioning Indexes 582
Local Indexes vs. Global Indexes 583
Local Indexes 584
Global Indexes 590
Partitioning and Performance, Revisited 606
Auditing and Segment Space Compression 612
Summary 614

CHAPTER 14 Parallel Execution 615

When to Use Parallel Execution 616
A Parallel Processing Analogy 617
Parallel Query 618
Parallel DML 624
Parallel DDL 627
Parallel DDL and Data Loading Using External Tables 628
Parallel DDL and Extent Trimming 630
Parallel Recovery 639
Procedural Parallelism 639
Parallel Pipelined Functions 640
Do-It-Yourself Parallelism 643
Summary 648

CHAPTER 15 Data Loading and Unloading 649

SQL*Loader 649
Loading Data with SQLLDR FAQs 653
SQLLDR Caveats 679
SQLLDR Summary 680
External Tables 680
Setting Up External Tables 681
Dealing with Errors 687
Using an External Table to Load Different Files 690
Multiuser Issues 691
External Tables Summary 692
Flat File Unload 692
Data Pump Unload 701
Summary 703

INDEX 705